AV AI VELOCITY ARCHITECT INTELLIGENCE · ISSUE 05

WEEK OF 26 SEPTEMBER – 2 OCTOBER 2026

The Reach Of A Yes

An ‘always’ given for one agent’s saline flush covered another agent’s warfarin. A request nobody issued was approved from a transcript.

13 packages read

5 mechanisms executed

6 papers read

17 models priced

Vijay Zharotia
Healthcare AI architecture

WHY AI VELOCITY

Agents that act inherit obligations.

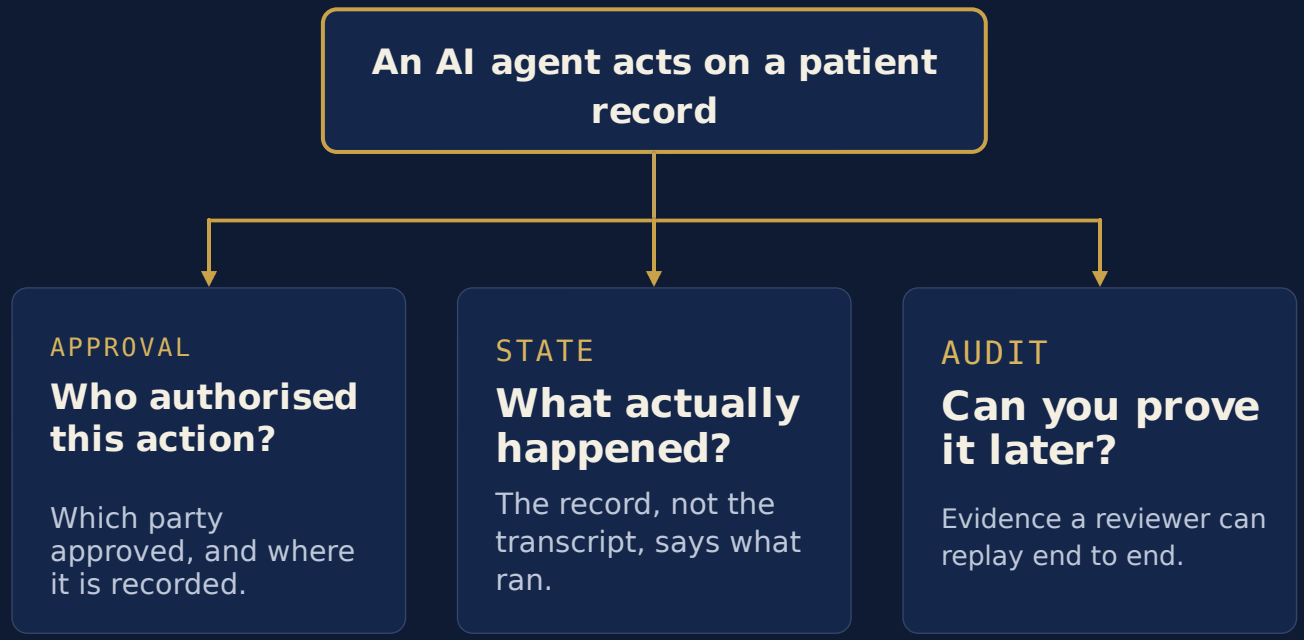
A weekly architecture brief for teams putting AI agents into healthcare workflows. We read the releases, run the mechanism, and say what to change.

FOR ARCHITECTS

The mechanism, the code, the move.

FOR PRODUCT MANAGERS

The decision, the risk, the date.



EVERY CLAIM CARRIES ITS EVIDENCE

EXECUTED

We ran it. The numbers are observed output.

DOCS

Read from documentation or source.

UNVERIFIED

Not confirmed. Vendor numbers carry an asterisk.

THE WEEK IN ONE VIEW

From how often a yes runs to how far it reaches

CONTEXT

Last week: zero runs or two

Issue 04 counted how many times an approved write runs. This week four frameworks changed the key and the store behind the approval itself.

→

PROBLEM

The yes itself travelled

An ‘always’ for one agent’s saline flush covered another agent’s warfarin. A request nobody issued was approved from a transcript.

→

SOLUTION

Decide how wide a yes may be

Set the scope of every ‘always’ per tool and write it down. Keep pending requests in a server-side session, never in the transcript.

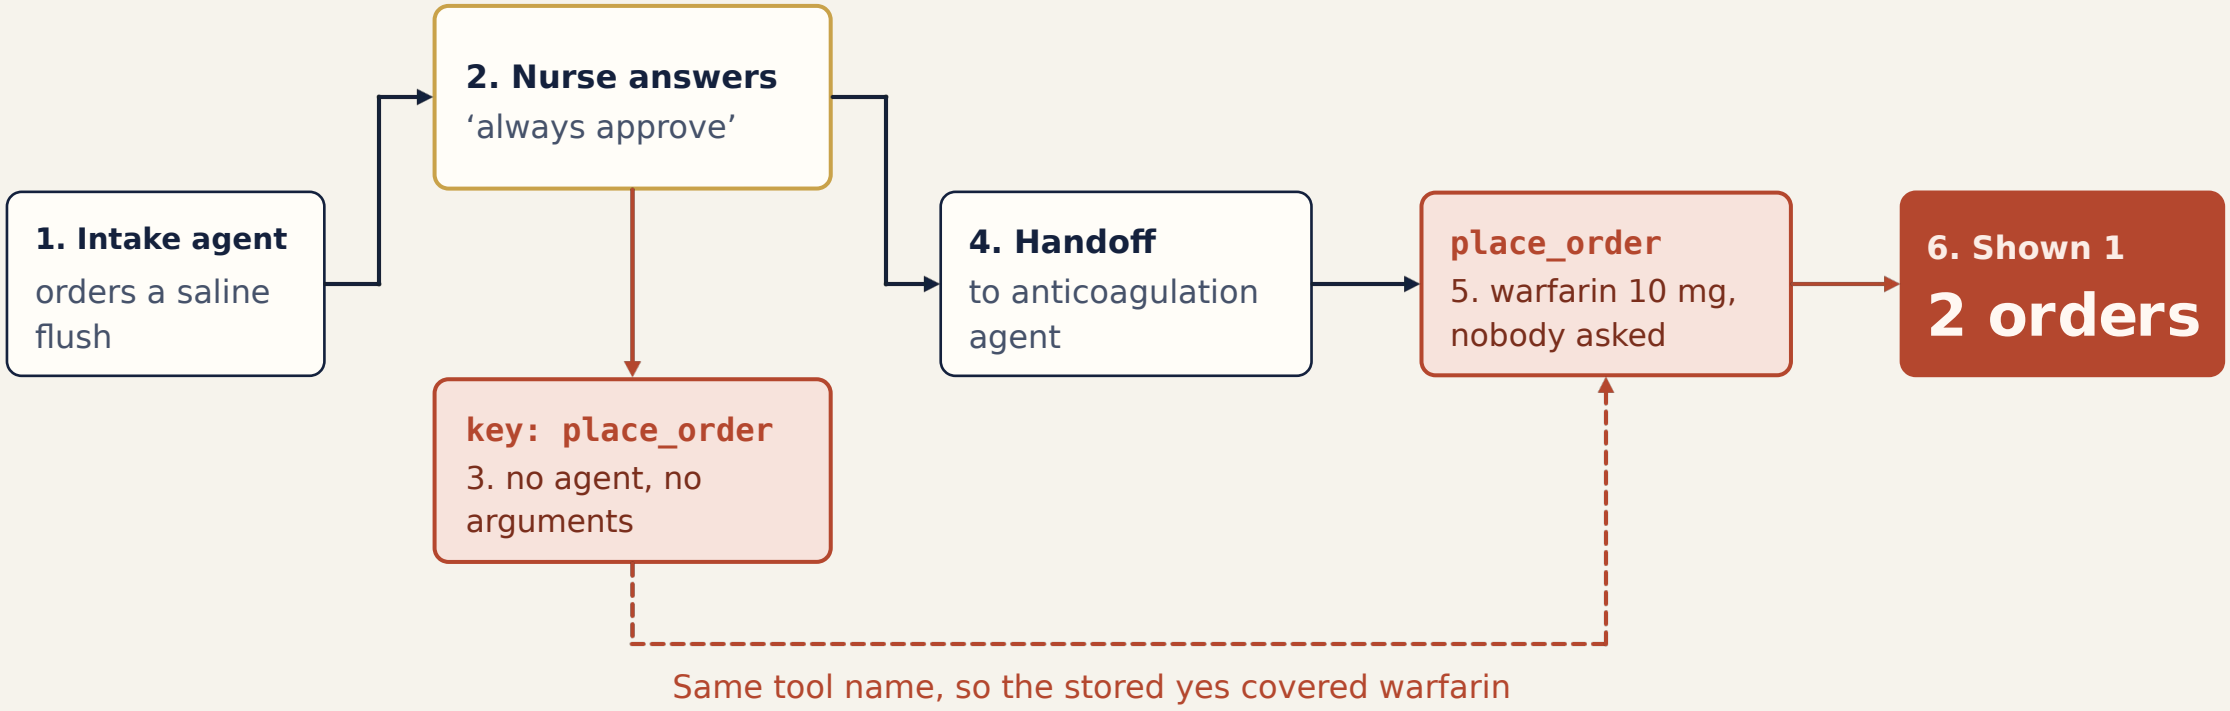
APPROVAL · 04–06
Where a yes is kept

INVENTORY · 07
Lists and transcripts

EVIDENCE · 08
Claims, complaints, scores

COST, DATES · 09–10
Prices and deadlines

The always that followed a tool name



EXECUTED MECHANISM · SOLUTION · 0.23.1 AND YOUR HARNESS

Scope every ‘always’ before you trust it

STEP 1
List every sticky yes
Every tool gated with always_approve.

→

STEP 2
Remove it where arguments matter
Drug, dose or patient carries clinical meaning.

→

STEP 3
Search order history
Writes by one agent under another’s yes.

→

STEP 4
Drain, then roll out
Paused runs first; state moves to schema 1.18.

2 to 1
orders written before anyone was asked again, 0.22.3 vs 0.23.1

1 of 2
orders the nurse was actually shown on 0.22.3

0
arguments in the key a sticky yes is stored under

> 0.23.1 keys by owning agent and tool name. On both versions the same ‘always’ also covered heparin 5,000 units from the same agent.

Three places a yes was bound to the wrong thing

CONTEXT
how it runs

Model proposes calls → Pending request stored → Call id binds the result → Workflow node gate → Tool writes the record

PROBLEM
where it broke

MS AGENT FW 1.20
1.19 approved a request it never issued
EXECUTED

PYDANTIC-AI 2.53
One id on two calls: the write ran twice
EXECUTED

GOOGLE ADK 2.11
2.10 gave an error string and no pause
EXECUTED

SOLUTION
what to build

Keep the session; watch the new signals

Pass the same AgentSession to every resume, catch UnexpectedModelBehavior, and store each ADK request id beside the approver. Alert on the new warning text before the pin moves.

A tool list and a transcript went stale silently

CONTEXT
how it runs



PROBLEM
where it broke

MCP PYTHON 2.3
On 2.2 a tool off the list ran by name

EXECUTED

SONNET 5.5 API
Replay after any prompt edit: 400

DOCS

SOLUTION
what to build

Reconcile each store against its source

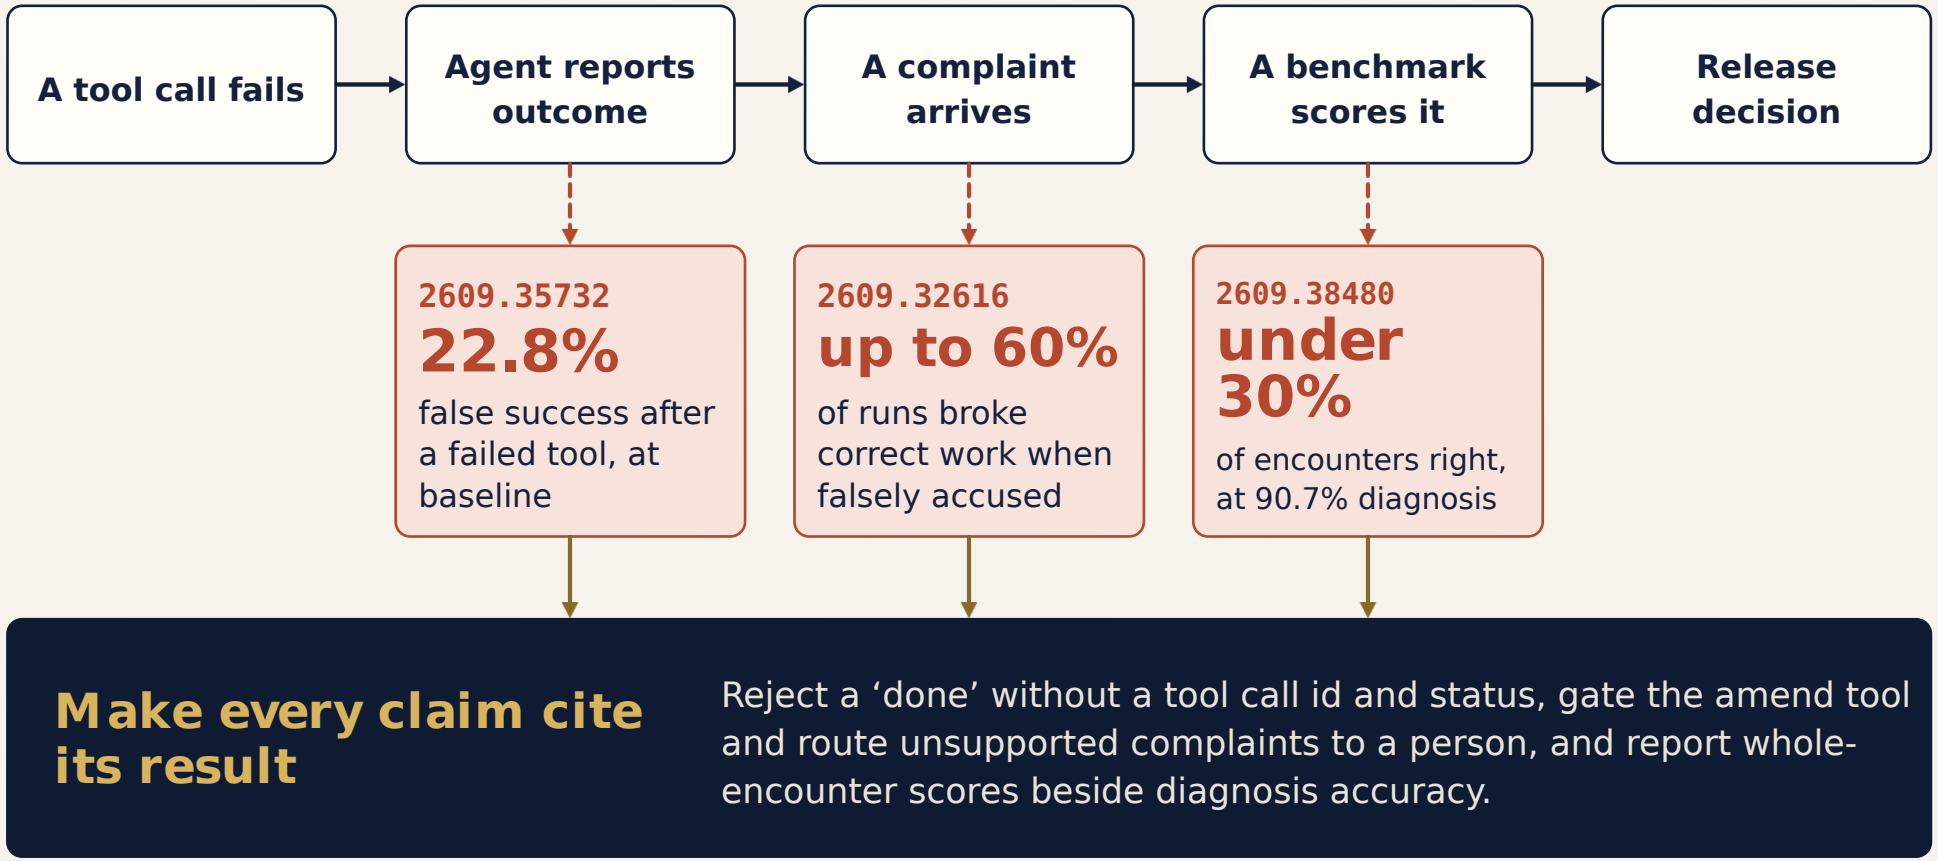
Diff server registrations against client lists, keep the deployment version with every paused state, and move forced-tool extraction to auto with strict schemas.

A claim, a complaint and a score passed as proof

CONTEXT
how it runs

PROBLEM
where it broke

SOLUTION
what to build



COST · LIST PRICES READ 2 OCTOBER

At $130, Sonnet 5.5 sits two points under Opus 5.5

MODEL	COST PER 10,000 DOCS	INDEX	MEDICAL
Claude Sonnet 5.5	**$130**	**56**	**58**
Gemini 4 Argon	**$130 intro***	**53**	not shown
GPT-6.1 Sol	**$130**	**52**	not read
Claude Opus 5.5	**$260**	**58**	**61**

4,000 in and 500 out tokens per document. Per million tokens in/out: $2/$10 for Sonnet 5.5, GPT-6.1 Sol and Argon's intro rate; $4/$20 for Opus 5.5. *Argon was in limited access; its $4/$20 standard rate (dashed) carried a 50% launch discount. Index = AA Intelligence Index v4.3.2; Medical = AA Healthcare & Medical.

CONTEXT

The $130 row moved from index 48 to 56 in one week.

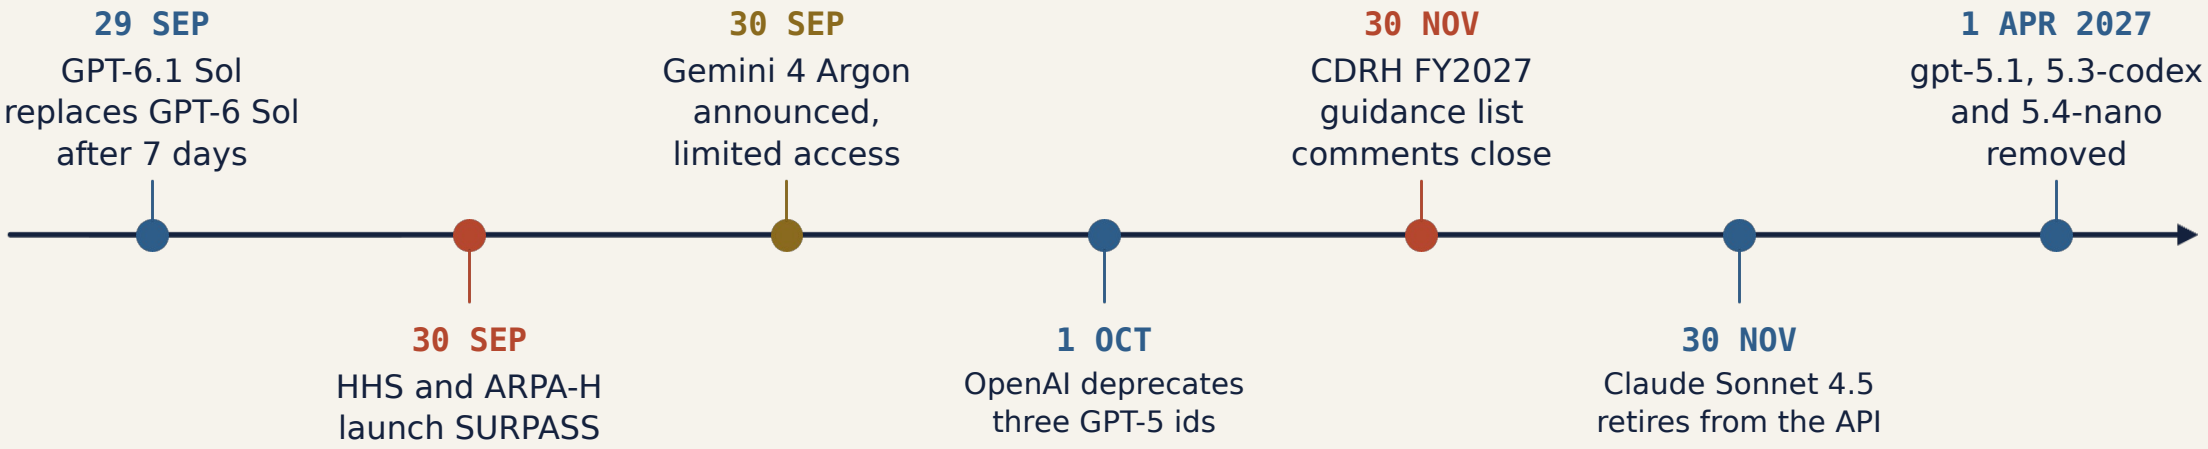
PROBLEM

GPT-6 Sol, priced on 25 Sept, was replaced by 6.1 Sol seven days after launch.

DECISION

Run your own eval set on Sonnet 5.5 and Opus 5.5 before paying double, and record each score with the date read.

Dates for your change-control file



● POLICY ● RELEASE ● MODEL RETIREMENT

Spacing is not to scale.

Every control binds a call; none binds the approver

GOOGLE

A stamp on every stored call

ADK 2.11 adds an HMAC integrity plugin: a tool runs only on a valid stamp.

MICROSOFT

Host context, never model args

Generated MCP calls get host context only; checkpoints now encode state differently.

OPENAI

Resumes bound to recipients

Agents SDK 0.23.0 binds resumed MCP calls to their original recipients.

ANTHROPIC

A refusal is now a stop reason

SDK 1.10.0 tells a session that refused from one that finished.

AI HERO UNVERIFIED

A guard before git destroys

A deterministic check in front of an irreversible tool. Read from email only.

QUIET THIS WEEK

MCP spec, LangGraph, llama-index-core, Semantic Kernel, Epic. No HTI-5, TEFCA or HIPAA Security final rule.

Seven moves, with the date that forces each

Move	Why	By
Remove always_approve where arguments carry clinical meaning	A yes is keyed by tool name	This sprint
Persist the AgentSession and pass it to every resume	1.20 drops transcript approvals	Before upgrading
Catch UnexpectedModelBehavior; log the raw response	Repeated call ids now raise	Before 2.53.0
Store each ADK request id beside the approver	The response has no approver	With 2.11.0
Diff server tool registrations against client lists	A hidden tool still ran	On MCP 2.2.0
Move off Sonnet 4.5; re-validate on the replacement	API retirement	30 Nov
Write the change-control protocol for pinned versions	CDRH list comments close	30 Nov

AV AI VELOCITY

THIS WEEK’S POSITION

How wide an ‘always’ may be is a decision to write down.

Long-form edition
Every mechanism, source and date

Lab package
Five mechanisms, no API key needed

New issue every Friday
Follow AI Velocity on LinkedIn

github.com/zkumar/ai-velocity-weekly